

# **OXIEDO: The Architecture That Unlocks Locked Data**

The most valuable data in the world cannot be trained on.  
ORMAS produces the account that  
changes that, as the model learns it.

**383 Controlled Experiments** • Reviewer, NeurIPS 2026 Workshop • Submitted to AAAI 2027

# | One Decision in 1986. Three Problems Everyone Funds Separately.

---

## A Single Global Error Signal Makes Fault Localization Mathematically Impossible — Every Downstream Failure Follows From It

---

### NOTHING CAN BE LOCATED

One error signal spreads across every parameter at once.

No unit can assess its own state.

The loss curve collapses and nothing says which node died.

### NOTHING IS RETAINED

New tasks overwrite prior representations indiscriminately.

ResNet-18 keeps 47.3% after a sequential task shift.

*Every regime change costs what the model knew.*

### NOTHING IS REPAIRED

Repair requires knowing what is broken.

Corrupted data is learned silently and confidently.

Rollback or full retrain are the only recourse.

*One cause: the chain rule multiplies every unit's contribution through every other unit's. Unchanged since 1986.*

# | Kill a Converged Layer. One Architecture Comes Back.

---

## Bounding Each Node's Local Gradient Chain Produces Attribution as a Physical Byproduct — and Attribution Makes Repair Possible

---

### STANDARD BACKPROPAGATION

One global error signal, entangled across all parameters.

- No per-node visibility.
- Silent collapse on node death.
- No diagnosis, no targeted repair.

Feature layer zeroed at epoch 100:  
**10.0%** — chance, every seed, zero variance, permanent.

### ORMAS THREE-SIGNAL TRAINING

Each node reads through a **4,416-parameter** bottleneck — too small to memorise noise.

- Scores its own health every step.
- Fires at  **$2.5\sigma$**  below its own history.
- Diagnoses 7 modes, repairs mid-run.

Same attack: **80.3%** after 85 targeted corrections (**+70.3pp**).

*Local stability characterization from Input-to-State Stability (Sontag, 1989) — not a global convergence proof, not yet peer-reviewed.*

# | The Industry Is Moving From Inference to Training on Private Data

---

## Three Independent Pressures Converge on the Same Missing Layer — None of Them Existed Two Years Ago

---

### REGULATORS SHIFTED TO ARCHITECTURE

The FDA rejects medical models for opacity **regardless of accuracy**.

The EU AI Act attaches liability to systems with no audit trail.

*The requirement is no longer performance. It is provenance.*

### RENTED INFERENCE STOPPED SCALING

Per-call cost compounds linearly with usage.

At institutional volume the only exit is your own infrastructure.

**Which means training on your own data — the thing nobody can do safely.**

### MODELS DEGRADE INVISIBLY

Every regime change, patient cohort and fraud pattern erodes what the model knew.

No detection exists at the training layer.

*It surfaces as a miscalculation, not an alert.*

*The regime is shifting from inference on public data to training on proprietary institutional data.*

# | Five Sectors. One Question. Counted Up, Not Down.

These Five Are Where the Missing Account Already Costs Somebody Money, in Figures That Can Be Checked

## THE QUALIFYING BUYER

Three properties at once, or not a customer:

1. Data that cannot legally leave.
2. An active plan to train on it.
3. Somebody who will ask what the model learned.

*The constraint keeps the count small and the contracts large.*

## WHAT THE GAP COSTS TODAY

**AI Training** — \$16K–\$24K per recovery from a silent failure.

**Medical AI** — 12–36 months added to market entry.

**Regulated Finance** — no component-level audit, no validation. Programme written off.

**Data Obligation** — \$100K–\$500K per deletion request.

**Defense** — “we cannot say how” ends the deployment.

*Published third-party figures, one per sector. No buyer interviews yet — the primary open risk, stated rather than assumed away.*

# | One Product. Three Things Arrive With the Model.

---

Training Is Unchanged — Data, Objective and Resulting Model Are Identical. What Changes Is the Account That Exists Alongside It

---

## 01 — THE MODEL

*Deploys exactly as it would have*

Same architecture, same inference cost, no change to the serving path.

Nothing is traded away for the accountability. That is the whole point.

## 02 — THE NAMED PARTS

*What it learned, per source*

Each source's contribution held as a discrete addressable object, bound to the record of the data that justified it.

Removable on instruction, without retraining.

## 03 — THE RECORD

*Every change training made*

Which component, at which step, under which diagnosis, by how much, inside what declared limit.

Signed, and diffable against the last approved version.

*Seven features of one architecture, not seven products. Nothing is priced per feature. [anonymous.4open.science/r/ormas-EB73](https://www.reddit.com/r/openai/comments/10m7k7k/anonymous_4open_science_rormas_eb73_all_383_runs_reproduce_with_one_command/) — all 383 runs reproduce with one command.*

# I Registered a Company to Try to Disprove My Own Architecture

## Seven Escalating Tests Over Twelve Months, Each Designed to Fail in Public — All Seven Held

### RUNGS 1–4 REMOVE IT

1. **Origination** — cold start, no brand, no ads. Customers arrived, 100% LLM-referred.
2. **Unit economics** — \$0.0043 per 12-asset suite against \$50–\$150 human.
3. **Causal isolation** — system stripped, everything else identical. –91% output.
4. **Collapse shape** — –100% new customers, every channel on the same day.

### RUNGS 5–7 PUT IT BACK

5. **Reversibility** — re-injected, nothing else changed. Recovered.
  6. **Magnitude** — +1,300%, overshooting 3.3× on accumulated memory alone.
  7. **Ceiling** — \$4,386 order closed at \$0.00 CAC, no human in the transaction.
- It stops there because I stopped it** — no jurisdiction assigns liability to an agent.

*This tests the orchestration layer, not ORMAS — prior work, now closed, with its own separate record. Twelve-month revenue: \$6,691.68 across 79+ orders, capped at research grade deliberately.*

# | Publish the Mechanism. License the Accountability.

---

## Defensibility Never Rested on Secrecy — It Rests on Calibration, Regulatory Acceptance, and a Record a Reimplementation Does Not Come With

---

### PUBLISHED — FREE

The architecture and every experiment behind it release in full on publication.

A result nobody can check is worth nothing to a regulator.

*Written into the licence, not promised.*

### LICENSED — REVENUE

Production deployment, on the customer's own infrastructure.

One integration, one contract, nothing priced per feature.

Zero data leaves the perimeter, so margin is software-grade.

### WHAT IS BOUGHT

Not the architecture — that is free from release day.

A calibrated bound, and a counterparty who will stand behind it in a filing.

*A reimplementation arrives with none of the record.*

*Zero contracts signed, zero customer conversations held. Stated here rather than found in week three of diligence.*

# | The Regulator Creates the Demand. A Person Has to Close It.

---

## The Requirement Is Written Down Before We Arrive — What Is Missing Is Someone Who Has Survived a Procurement Review

---

### WHY THEY COME TO US

- Regulators now reject models on architecture, not accuracy. The requirement exists before the first meeting.
- The audit trail is produced during training, not reconstructed afterwards.
- Nothing leaves the client's infrastructure — which removes the objection that kills every hosted alternative.

### THE ONE GAP

A regulated institution buys from a person it can hold accountable in the room.

**No architecture sits in a compliance review and answers for itself.**

**That is the co-founder** — someone who has sold into a fund, a hospital or an insurer and won.

*Design partner → single-model deployment → organisation-wide licence. Regulated cycles run 6–18 months.*

# Everyone Is Competing One Layer Above the Constraint

Orchestration Is Not the Blocker — the Blocker Is That You Cannot Safely Train on the Data

|                           | Per-node<br>attribution | Stability<br>analysis | Autonomous<br>repair | Audit trail<br>at train time | Their<br>metal |
|---------------------------|-------------------------|-----------------------|----------------------|------------------------------|----------------|
| Agent frameworks          | ×                       | ×                     | ×                    | ×                            | ×              |
| Enterprise data platforms | ×                       | ×                     | ×                    | ~                            | ~              |
| Frontier model vendors    | ×                       | ×                     | ×                    | ×                            | ×              |
| Data labelling vendors    | ×                       | ×                     | ×                    | ×                            | ×              |
| Academic research         | ~                       | ~                     | ~                    | ×                            | —              |
| <b>This</b>               | ✓                       | ✓                     | ✓                    | ✓                            | ✓              |

The moat is not secrecy — the work publishes on release. It is that getting here ran through a live company broken on its own data, then control theory, then two years of experiments. That order does not compress.

# | Everything Technical Is Built. One Thing Is Not.

---

## Three Years Solo Produced the Architecture, the System and the Causal Proof — the Gap Is a Person, Not a Capability

---

### FOUNDER

**Eighteen.** No university, no lab, no advisor, no external funding.

A self-correcting architecture with a formal stability result, derived at 17.

Programme committee reviewer, **NeurIPS 2026** workshop.

Every experiment ran on one GPU, bought at 15 with money from the first thing he sold.

### CO-FOUNDER SOUGHT

**One specific profile:**

Has sold software inside a fund, a hospital or an insurer.

Understands model risk well enough to win an argument with a compliance officer.

**Not an engineer. Not a generalist.**

*Equal equity, four-year vest.*

*The honest gap: never having worked inside a team. [raadh.me](https://raadh.me) · [zenodo.org/records/21730363](https://zenodo.org/records/21730363)*

# | From a Training Layer to the Substrate Regulated AI Runs On

---

Every Phase Removes One Dependency — Ending With a Stack That Rents Intelligence From No One

---

## PHASE 1 — NOW

Finish the research and carry per-node health monitoring onto attention heads.

Architecture-agnostic by design, so this is a compute problem, not a research one.

First design partners.

## PHASE 2 — 12-24 MONTHS

First licensed deployments inside institutions that cannot move their data.

Each accumulates institutional memory that cannot be exported.

Switching cost compounds with tenure.

## PHASE 3 — YEARS 3-5

A model pre-trained on the self-correcting architecture from scratch.

The external dependency disappears; cost collapses to hardware.

A position no competitor reaches by adding a feature.

**The network learns to stop needing correction. That is the thesis, and it is measured.**

*Correction frequency decays  $4.2 \rightarrow 0.05$  per epoch — the system repairs itself until it no longer needs to.*

# | \$1.5M Pre-Seed. Sized to the Next Result, Not the Whole Plan.

## Eighteen Months, Six Milestones With Acceptance Criteria — the Compute Line Costed Against the Question Actually Open

### USE OF FUNDS

- \$650K Co-founder + research engineer
- \$250K 32–64 accelerators, two campaigns
- \$150K One custodian, non-exclusive
- \$100K Delaware, EU rep, filings
- \$200K 18 months of everything else
- \$150K Buffer, unallocated on purpose

### WHAT IT BUYS, IN ORDER

1. **A second and third person.** Key-person risk is total today.
2. **One result that is not CIFAR.** Bounded attribution on a transformer backbone, same adversarial suite.
3. **One custodian saying yes.** No regulated data has ever touched this.

*The frontier-scale campaign is a seed-round problem. It is not being asked for here.*

Before writing a cheque, spend \$50 instead: the remaining programme is falsifiable in under 30 GPU-hours, on one card, on your hardware.